

POCKET PRODUCER: SESSION-CONDITIONED RETRIEVAL AS USER-CONTROLLED PROPOSAL

Chuling Li

University of Nottingham

scycl14@nottingham.ac.uk

ABSTRACT

We present Pocket Producer, a prototype that operationalizes *session-conditioned continuation retrieval* inside the Audiotool DAW. The system reads a live session via the Nexus SDK, computes a session fingerprint from tempo and CLAP embeddings of up to eight region-referenced samples, and ranks Freesound Loop Dataset candidates with a five-signal logistic fusion. Candidates are surfaced for preview with per-signal evidence; insertion is a single reversible Nexus transaction; feedback is logged for future personalization. Offline evaluation on pack co-membership weak labels shows the fusion advantage is sensitive to label construction, representation space, and negative regime—motivating the design choice to default to a transparent rules baseline and treat the learned ranker as a selectable *proposal*, not an autonomous decision.

1. INTRODUCTION

When a musician returns to an unfinished loop-based project, the question is not “what sounds similar?” but “what would be useful as the next step?” Loop compatibility [1–4], stem coherence [5–7], session-conditioned generation [8], and reversible co-creation [9] all address parts of this problem. We define **continuation utility** as the property that a candidate helps the musician make progress in the current session; unlike similarity it is asymmetric and context-dependent.

We contribute: (1) a DAW-integrated prototype with session fingerprinting, multi-signal ranking, and a proposal-based interaction model where every recommendation is previewable and reversible; (2) a cold-start ranking recipe using Freesound pack co-membership as weak supervision; (3) empirical observations showing the proxy is too sensitive to label construction, embedding space, and negative regime to serve as an autonomous ranker.




Figure 1. Pocket Producer pipeline: the system proposes (steps 1–3); the musician decides (steps 4–6).

2. SYSTEM

2.1 Session Fingerprinting

The live Audiotool session (Figure 1) is read via the Nexus SDK, which exposes the project’s devices, regions, and referenced audio samples in real time. The system extracts the project tempo from the document metadata and collects up to eight unique audio samples referenced by regions in the arrangement. Each sample is embedded with CLAP [10]; the resulting vectors form the session’s timbral fingerprint. When no session audio is available (empty project or no sample-backed regions), the system falls back to non-audio signals only.

2.2 Multi-Signal Ranking

Five signals score each FSLD [11] candidate against the session fingerprint: (i) mean-pooled and (ii) max-pooled CLAP cosine similarity to session samples, (iii) tempo proximity, (iv) key compatibility, and (v) tag Jaccard overlap between candidate tags and tags from the session’s audio fragments. A logistic fusion trained with BPR loss [12] on Freesound pack co-membership weak labels (Section 3) combines these into a single score. A hand-tuned *rules* baseline using only tempo, key, and tag signals serves as the default ranker; the fusion is selectable and falls back to rules when no session audio is available.

2.3 Proposal Interface

The title’s “user-controlled proposal” is a deliberate design choice, not a fallback from poor accuracy. Pack co-membership is a proxy for timbral compatibility, not continuation utility; even a perfect ranker under this label would not capture arrangement context, creative intent, or



genre conventions. The interface therefore keeps the musician in the loop at every step.

Each candidate is surfaced for in-context preview with per-signal evidence—cosine score, tempo delta, key match, and overlapping tags—so the musician can assess *why* the system ranks it highly and judge whether those reasons apply to their current goal. An accepted candidate is inserted via a single Nexus transaction; *undo* reverts the insertion and retracts the candidate’s tags from the ranking context, so subsequent recommendations are not biased by the undone choice. Acceptance, rejection, and undo events are logged with session context for future personalization (not yet evaluated).

2.4 Deployed Model

The deployed ranker follows the same five-signal recipe but is a separate production model with four stated deltas: MS-CLAP embedding space (offline evaluation uses LAION-CLAP Music); source-weighted BPR (hard_similar $\times$ 2.0, mean over 5 seeds); key constant zero online (Audiotool documents lack key-signature metadata); tag context drawn from session fragments rather than pack siblings.

3. WEAK-LABEL FUSION

From the FSLD [11] we kept only sounds with a true Freesound pack (1,797 items, 605 packs), discarding items grouped only by uploader. Each eligible pack with at least three items yields one leave-one-out query per item, with same-pack siblings as positives and cross-pack items as negatives in three difficulty tiers (*easy*: tempo/key mismatch; *hard_tempo_key*: compatible tempo and key; *hard_similar*: highest cosine to the positive from another pack). A 121-source held-out split (seed 20260810) was committed to version control together with the evaluation protocol *before* any evaluation ran. Development uses only the 484 train sources; the 171 with ≥ 3 items form queries (981 examples), the rest contribute negatives only. LAION-CLAP Music [13] (512-dim) provides the embedding backbone.

On the train side (mean over 5 source-grouped splits), frozen cosine reaches 0.929 overall and 0.814 on hard_similar; five-signal fusion reaches 0.925 and 0.813. Under pack-only labels the fusion advantage is gone: -0.03 pp on hard_similar, and fusion trails cosine on hard_tempo_key by 2.1 pp on average (negative in all 5 splits). These development-set numbers are descriptive; inference is reserved for the held-out analysis (Section 4).

4. OBSERVATIONS

Four lines of evidence jointly motivate treating the fusion as a proposal rather than a decision. **Label sensitivity.** An earlier corpus treated same-uploader items as co-members (26% of examples, 949 sources); fusion appeared $+4.0$ pp over frozen cosine on hard_similar. Restricting to true pack co-membership erases the gain (Section 3). **Regime dependence.** Leave-one-signal-out ablation on the dev set

shows tempo helps on hard_similar (-0.96 pp when removed) but hurts on hard_tempo_key ($+0.68$ pp); signals switch sign across negative regimes, so a fused score averaging over them cannot be uniformly trusted. **Held-out uncertainty.** On a pre-committed 121-source split, the fusion—cosine advantage on hard_similar is $+1.5$ pp (bootstrap 95% CI $[-1.5, +4.7]$, crossing zero) under LAION-CLAP; under MS-CLAP with source-weighted BPR it is $+9.2$ pp (CI $[+4.3, +14.3]$)—a 7.7 pp swing evidencing representation sensitivity. **Pilot divergence.** Nine raters on 20 blind A/B pairs (Fleiss’ $\kappa = 0.294$) show a descriptive divergence between weak-label signal rankings and human consensus.

5. CONCLUSION

Pocket Producer operationalizes session-conditioned retrieval as a user-controlled proposal in a production DAW. It reads the live session via the Nexus SDK, ranks candidates with a multi-signal fusion, and surfaces per-signal evidence for preview; every insertion is reversible and every interaction is logged. Since offline evaluation shows the fusion advantage is sensitive to label construction, embedding space, and negative regime, the default ranker is a transparent rules baseline; the fusion serves as a selectable alternative. Evaluating whether preview and undo improve creative outcomes, and personalization from logged feedback, is future work.

- 145 **Acknowledgements.** The author thanks the Audiotool 197
 146 team for the Nexus SDK and the FSLD creators [11]. 198
 199
 200
 201
- ## 6. REFERENCES
- 147
- 148 [1] B.-Y. Chen, J. B. L. Smith, and Y.-H. Yang, “Neural 202
 149 loop combiner: Neural network models for assessing 203
 150 the compatibility of loops,” in *Proceedings of the 21st 204
 151 International Society for Music Information Retrieval 205
 152 Conference (ISMIR)*, 2020. 206
- 153 [2] S. Lattner, “SampleMatch: Drum sample retrieval by 207
 154 musical context,” in *Proceedings of the 23rd Interna- 208
 155 tional Society for Music Information Retrieval Confer- 209
 156 ence (ISMIR)*, 2022. 210
- 157 [3] X. Yi and M. Barthet, “A generative framework for 211
 158 composition-aware loop recommendation in music pro- 212
 159 duction: Drum2Bass use case,” in *Proceedings of the 213
 160 Sound and Music Computing Conference (SMC)*, 2024.
- 161 [4] S. Roy Chaudhuri, S. P. Madupu, K. T. Guru Sai, 214
 162 and V. Jamwal, “LoopMatcher: Proof-of-concept for 215
 163 AI-assisted music loop search,” in *Proceedings of the 216
 164 15th International Conference on Computational In- 217
 165 telligence in Music, Sound, Art, and Design (Evo- 218
 166 MUSART)*, 2026.
- 167 [5] R. Ciranni, G. Mariani, M. Mancusi, E. Postolache, 219
 168 G. Fabbro, E. Rodolà, and L. Cosmo, “COCOLA: 220
 169 Coherence-oriented contrastive learning of musical au- 221
 170 dio representations,” in *IEEE International Conference 222
 171 on Acoustics, Speech and Signal Processing (ICASSP)*, 223
 172 2025.
- 173 [6] J. Nistal, M. Pasini, C. Aouameur, M. Grachten, and 224
 174 S. Lattner, “Diff-A-Riff: Musical accompaniment co- 225
 175 creation via latent diffusion models,” in *Proceedings of 226
 176 the 25th International Society for Music Information 227
 177 Retrieval Conference (ISMIR)*, 2024.
- 178 [7] A. Riou, A. Gagnere, G. Hadjeres, S. Lattner, and 228
 179 G. Peeters, “Zero-shot musical stem retrieval with 229
 180 joint-embedding predictive architectures,” in *IEEE In- 230
 181 ternational Conference on Acoustics, Speech and Sig- 231
 182 nal Processing (ICASSP)*, 2025.
- 183 [8] Q. He, G. Xia, and Z. Wang, “TOMI: Transforming 232
 184 and organizing music ideas for multi-track composi- 233
 185 tions with full-song structure,” in *Proceedings of the 234
 186 26th International Society for Music Information Re- 235
 187 trieval Conference (ISMIR)*, 2025.
- 188 [9] A. C. Elkins, S. Singh, A. Kieback, S. Blankenship, 236
 189 U. P. Amadasun, and A. Chadha, “DAWZY: A new 237
 190 addition to AI powered “human in the loop” music co- 238
 191 creation,” in *NeurIPS Workshop on AI for Music*, 2025.
- 192 [10] B. Elizalde, S. Deshmukh, M. Al Ismail, and H. Wang, 239
 193 “CLAP: Learning audio concepts from natural lan- 240
 194 guage supervision,” in *IEEE International Conference 241
 195 on Acoustics, Speech and Signal Processing (ICASSP)*, 242
 196 2023.
- 197 [11] A. Ramires, F. Font, D. Bogdanov, J. B. L. Smith, Y.-H. 243
 198 Yang, J. Ching, B.-Y. Chen, Y.-K. Wu, W.-H. Hsu, and 244
 199 X. Serra, “The freesound loop dataset and annotation 245
 200 tool,” in *Proceedings of the 21st International Society 246
 201 for Music Information Retrieval Conference (ISMIR)*, 247
 202 2020.
- 203 [12] S. Rendle, C. Freudenthaler, Z. Gantner, and 248
 204 L. Schmidt-Thieme, “BPR: Bayesian personalized 249
 205 ranking from implicit feedback,” in *Proceedings of the 250
 206 25th Conference on Uncertainty in Artificial Intelli- 251
 207 gence (UAI)*, 2009.
- 208 [13] Y. Wu, K. Chen, T. Zhang, Y. Hui, T. Berg-Kirkpatrick, 252
 209 and S. Dubnov, “Large-scale contrastive language- 253
 210 audio pretraining with feature fusion and keyword- 254
 211 to-caption augmentation,” in *IEEE International Con- 255
 212 ference on Acoustics, Speech and Signal Processing 256
 213 (ICASSP)*, 2023.