

RE-ENTERING UNFINISHED MUSIC: SESSION-CONDITIONED RETRIEVAL OF PERSONAL AUDIO FRAGMENTS WITH REVERSIBLE DAW ACTIONS

Chuling Li

Independent Researcher

chuling.li.cs@gmail.com

ABSTRACT

Musicians accumulate large libraries of unfinished audio fragments yet get little support deciding what to try next when re-entering a project. We formulate *session-conditioned continuation retrieval*: given the live state of a DAW project, rank a musician’s own fragments by their utility for the next creative step. Pocket Producer implements this in Audiotool: a transparent fingerprint of the live project, ranking with frozen audio–text representations, and reversible single-transaction insertion. On 1,755 weak pairs from the Freesound Loop Dataset, frozen CLAP cosine reaches 0.883 pairwise accuracy and learned rerankers (up to 335K parameters) cannot beat it; a five-parameter fusion of cosine, tempo, key, and tag signals reaches 0.794 on the hardest negatives, +4.0 pp over cosine (0.754; one-sided sign test $p = 0.031$). Session-structure signals (instrument role gap, harmonic compatibility) show no gain on pack-membership labels, while an exploratory alignment of the same signals with external rater preferences (9 raters, 20 pairs, Fleiss’ $\kappa = 0.294$) points the opposite way—weak-label and human evaluations appear to reward different structure, motivating task-specific evaluation.

1. INTRODUCTION

When a musician returns to an unfinished loop-based project, the central question is not “what sounds similar?” but “what would be useful as the next step?” Existing work on loop compatibility [1–3] and stem coherence [4, 5] asks whether two elements can co-exist in a mix; our query is a *live session state*, and the candidates are the musician’s own creative memory, not a public catalog.

We define **continuation utility** as the property that a candidate fragment helps the musician make progress in the current session—subsuming tempo fit, timbral contrast, structural gap-filling, and stated intent. We contribute:

1. **Task:** session-conditioned ranking of personal fragments with a preregistered primary metric (pairwise



Figure 1. Pocket Producer retrieval loop: live session fingerprint $\rightarrow$ ranking $\rightarrow$ reversible insertion $\rightarrow$ preference feedback.

accuracy over human preferences).

2. **System:** a working prototype in a real DAW (Audiotool) with live document reading, reversible insertion, and explainable evidence.
3. **Protocol:** a backbone bake-off, a baseline comparison (learned rerankers lose to frozen cosine; multi-signal fusion gains +4.0 pp on hard negatives), and a blind A/B evaluation with 9 external raters.

2. POCKET PRODUCER

Figure 1 shows the loop; the system reads and writes Audiotool projects through the Nexus SDK.

2.1 Session Fingerprint

The live Nexus document is mapped deterministically to a `SessionFingerprint` (tempo, track counts and types, playhead, user-editable intent); only this summary leaves the browser, and missing fields are never imputed. Session audio regions are encoded with MS-CLAP 2023 [6] into a mean session embedding (CLAP’s joint audio–text space also serves text intent); per-region instrument roles and 12-dim CQT chroma are extracted for the session-structure signals below.

2.2 Continuation Utility Ranking

Candidate fragments are scored in two stages. First, per-space retrieval runs independently over four signal spaces: (1) structured metadata (tempo/key/tag matching), (2) frozen CLAP embeddings (cosine similarity), (3) *role-gap fill*—a zero-shot instrument probe boosts candidates whose predicted role fills a gap in the session, and (4) *harmonic compatibility*—transpose-invariant cross-correlation of CQT chroma vectors scores tonal fit



Table 1. Backbone bake-off on 2,558 FSLD items. R@10 = same-source Recall@10; Genre/Role = linear-probe accuracy.

Backbone	Dims	R@10	Genre	Role	Bytes/frag
laion-clap-music	512	0.713	0.415	0.556	2,048
msclap-2023	1024	0.692	0.384	0.533	4,096

Table 2. Baseline comparison on FSLD weak pairs (LAION-CLAP 512-dim, 5 seeds). Hard-sim = hard_similar negatives; Hard-tk = hard_tempo_key. Bold = best per column.

Model	Params	Overall	Hard-sim	Hard-tk
cosine-mean	0	0.883	0.754	0.911
deepsets-v1	164K	0.735	—	—
pocketrank-v1	335K	0.742	—	—
cosine+session	2	0.882	0.749	0.911
fusion (5-signal)	5	0.894	0.794	0.897
fusion (7-signal)	7	0.892	0.792	0.894

across all 12 semitone shifts. Results are fused with Reciprocal Rank Fusion [7]; a learned reranker can be applied on top, and a rule baseline with structured evidence codes is the always-available fallback.

2.3 Reversible Audiotool Interaction

A recommendation can be previewed and inserted in a single Nexus transaction; an `InsertReceipt` records created entity IDs so undo restores the exact prior state (verified with the official WASM validator); interactions are logged as preference feedback.

3. PRELIMINARY EVALUATION

3.1 Backbone Bake-off

On 2,558 items from the Freesound Loop Dataset (FSLD) [8], LAION-CLAP Music [9] (512-dim) outperforms MS-CLAP 2023 (1024-dim) on all three probes (Table 1) at half the storage. We deploy MS-CLAP 2023 because its training-data license audit is complete—quality and eligibility are separate decisions.

3.2 Baseline Comparison and Multi-Signal Fusion

We constructed 1,755 leave-one-out pairs from FSLD, using same-pack membership as a weak proxy for continuation utility, with three negative types: *easy* (tempo or key incompatible), *hard_tempo_key* (tempo and key both compatible), and *hard_similar* (highest embedding cosine from another pack). Splits are grouped by source pack; Table 2 reports 5-seed mean validation pairwise accuracy.

Frozen cosine similarity (0.883) is already a strong baseline: two learned rerankers with up to 335K parameters trained with BPR loss [10] score below it (0.735, 0.742), overfitting to noise in the weak labels. However, a five-parameter logistic fusion of audio cosine (mean and max pooling), tempo match, key match, and tag Jaccard overlap—trained with the same BPR loss—reaches **0.794**

on *hard_similar* negatives, a +4.0 pp gain over cosine alone (0.754); the gain holds in all 5 seeds (one-sided sign test $p = 0.031$). *Hard_similar* negatives are by construction cosine’s failure mode (mined as highest-cosine from another pack): the +4.0 pp measures what complementary signals recover there—the overall gain is +1.1 pp, holding in all 5 paired splits. On *hard_tempo_key* negatives, where tempo and key features are neutral by construction, fusion (0.897) stays close to cosine (0.911), ruling out label leakage.

Adding the session-structure signals (Section 2.2) does not improve the fusion: 7-signal results (0.892 overall, 0.792 *hard_similar*) are indistinguishable from 5-signal, and *cosine+session* (0.749) sits slightly below cosine alone—pack membership reflects recording context, not whether a candidate fills a compositional gap.

3.3 Author Pilot Evaluation

Under a preregistered protocol (frozen 2026-07-30), the author labeled 60 blind A/B pairs; 42 received a preference, 18 “neither” (30%). No ranker exceeds 0.571 pairwise accuracy (chance = 0.5), and cosine-agreement is 57.1% (24/42)—statistically indistinguishable from chance (Wilson 95% CI [0.42, 0.71]). The role-gap signal fires on only 7 of 42 pairs but agrees in 5 of 7. Author labels are excluded per the preregistration.

3.4 External Blind Evaluation

Nine external raters labeled the same 20 blind A/B pairs (180 labels). Fleiss’ $\kappa = 0.294$ indicates fair inter-rater agreement—typical for subjective music-preference tasks [11]—with 10 of 20 pairs reaching $\geq 7/9$ majority. Most pairs elicit a clear preference (only 8.9% “neither” versus the pilot’s 30%), suggesting context information aids decision-making.

In a post-hoc exploratory analysis (not preregistered; five signals compared), we align each signal with the per-pair human majority. On rater-consensus pairs ($\geq 7/9$), harmonic compatibility matches the majority in 7 of 9 cases, CLAP cosine in 6, max-pooled cosine in 4, and tempo—the strongest theory signal on weak labels—in 3; the sparse role-gap signal is correct on all 3 pairs where it fires, echoing the pilot’s 5/7. No single difference is significant at $n = 9$, but the direction is consistent: the signals that win on weak labels (Table 2) lose ground on human preferences and vice versa, suggesting the two evaluations reward different structure and motivating a preregistered replication.

4. DISCUSSION AND LIMITATIONS

Weak labels $\neq$ ground truth. Weak labels compare signals at scale but cannot establish continuation utility—and can invert the human signal ranking (Section 3.4). The fusion gain replicates on the deployed MS-CLAP space (+5.2 pp *hard_similar*, 5/5 seeds).

Scope and privacy. We measure perceived relevance, not creative quality; session audio is never stored and OAuth tokens never leave the browser.

155 **Acknowledgements.** The author thanks the Audiotool 206
 156 team for the Nexus SDK and the FSLD creators [8]. 207
 157 **AI Usage Statement.** Claude Code (Anthropic) was used 208
 158 for engineering implementation and documentation draft- 209
 159 ing. All research claims, experimental design, analysis, ci- 210
 160 tations, and figures are the sole responsibility of the author, 211
 161 who has verified each element against the linked artifacts 212
 162 and source publications. 213

163 5. REFERENCES 214

- 164 [1] B.-Y. Chen, J. B. L. Smith, and Y.-H. Yang, “Neural 216
 165 loop combiner: Neural network models for assessing 217
 166 the compatibility of loops,” in *Proceedings of the 21st*
 167 *International Society for Music Information Retrieval*
 168 *Conference (ISMIR)*, 2020.
- 169 [2] S. Lattner, “SampleMatch: Drum sample retrieval by
 170 musical context,” in *Proceedings of the 23rd Interna-*
 171 *tional Society for Music Information Retrieval Confer-*
 172 *ence (ISMIR)*, 2022.
- 173 [3] X. Yi and M. Barthelet, “A generative framework for
 174 composition-aware loop recommendation in music
 175 production: Drum2Bass use case,” in *Proceedings of*
 176 *the Sound and Music Computing Conference (SMC)*,
 177 2024.
- 178 [4] R. Ciranni, E. Postolache, G. Mariani, M. Mancusi,
 179 L. Cosmo, and E. Rodolà, “COCOLA: Coherence-
 180 oriented contrastive learning of musical audio repre-
 181 sentations,” *arXiv preprint arXiv:2404.16969*, 2024.
- 182 [5] J. Nistal, M. Pasini, C. Aouameur, M. Grachten, and
 183 S. Lattner, “Diff-A-Riff: Musical accompaniment co-
 184 creation via latent diffusion models,” in *Proceedings of*
 185 *the 25th International Society for Music Information*
 186 *Retrieval Conference (ISMIR)*, 2024.
- 187 [6] B. Elizalde, S. Deshmukh, M. Al Ismail, and H. Wang,
 188 “CLAP: Learning audio concepts from natural lan-
 189 guage supervision,” in *IEEE International Conference*
 190 *on Acoustics, Speech and Signal Processing (ICASSP)*,
 191 2023.
- 192 [7] G. V. Cormack, C. L. A. Clarke, and S. Büttcher, “Re-
 193 ciprocal rank fusion outperforms Condorcet and indi-
 194 vidual rank learning methods,” in *Proceedings of the*
 195 *32nd Annual International ACM SIGIR Conference*,
 196 2009.
- 197 [8] A. Ramires, F. Font, D. Bogdanov, J. B. L. Smith, Y.-H.
 198 Yang, J. Ching, B.-Y. Chen, Y.-K. Wu, W.-H. Hsu, and
 199 X. Serra, “The freesound loop dataset and annotation
 200 tool,” in *Proceedings of the 21st International Society*
 201 *for Music Information Retrieval Conference (ISMIR)*,
 202 2020.
- 203 [9] Y. Wu, K. Chen, T. Zhang, Y. Hui, T. Berg-Kirkpatrick,
 204 and S. Dubnov, “Large-scale contrastive language-
 205 audio pretraining with feature fusion and keyword-

to-caption augmentation,” in *IEEE International Con-*
ference on Acoustics, Speech and Signal Processing
(ICASSP), 2023.

- [10] S. Rendle, C. Freudenthaler, Z. Gantner, and
 L. Schmidt-Thieme, “BPR: Bayesian personalized
 ranking from implicit feedback,” in *Proceedings of the*
25th Conference on Uncertainty in Artificial Intelli-
gence (UAI), 2009.

- [11] M. Schedl, A. Flexer, and J. Urbano, “The neglected
 user in music information retrieval research,” *Journal*
of Intelligent Information Systems, vol. 41, no. 3, pp.
 523–539, 2013.